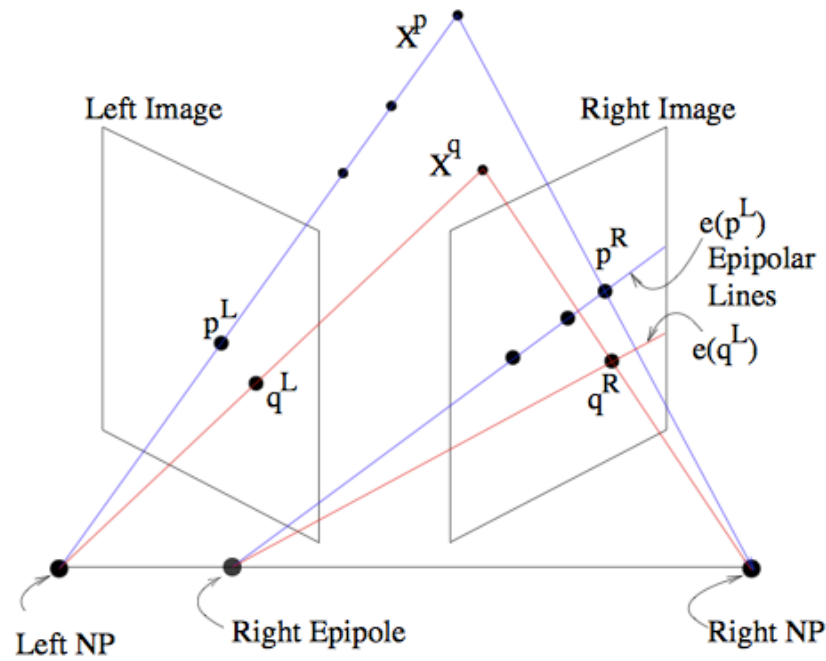


Projection, stereo images

X^p and X^q are the physical location of objects.

NP are the nadir points of the cameras.



The projection of the left nadir, NP is seen as a the right epipole with w.r.t. the right image. The epipoles may or may not be within the border of the images.

The projection of X^p to left nadir, NP is seen as an epipole line in the right image. If more than one epipole line is present in the right image, they converge at the right epipole (which might not be within the boundaries of the image).

Once the points in the left image, X_L are matched with points in the right image, X_R , if there are at least 7 points, one can determine the “bifocal tensor”, a.k.a. “fundamental matrix, relating the points in the 2 images using a 3x3 matrix of rank 2.

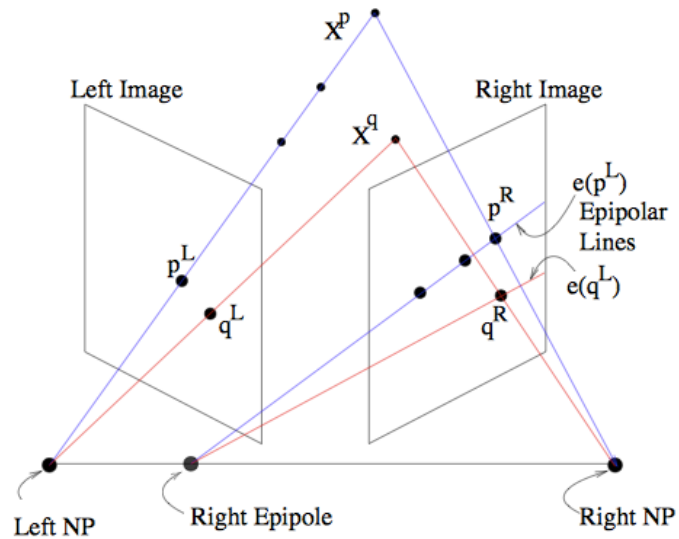
<http://www.cs.toronto.edu/~jepson/csc420/notes/epiPolarGeom.pdf>

(note, the image is posted on an educational site and copied here without following up on permissions. Any further use of the image should follow up on the origins and permissions.)

$(X_L)^T * F * X_R = 0$ for any pair or points in the images.



Note: the “Essential matrix” is a matrix used if the camera details are known. The “bifocal tensor”, a.k.a. “fundamental matrix” does not need camera details.

Projection, stereo and panoramic images



<http://www.cs.toronto.edu/~jepson/csc420/notes/epiPolarGeom.pdf>
 (note, the image is posted on an educational site and copied here without following up on permissions. Any further use of the image should follow up on the origins and permissions.)

panoramic images here are from Brown & Lowe 2003


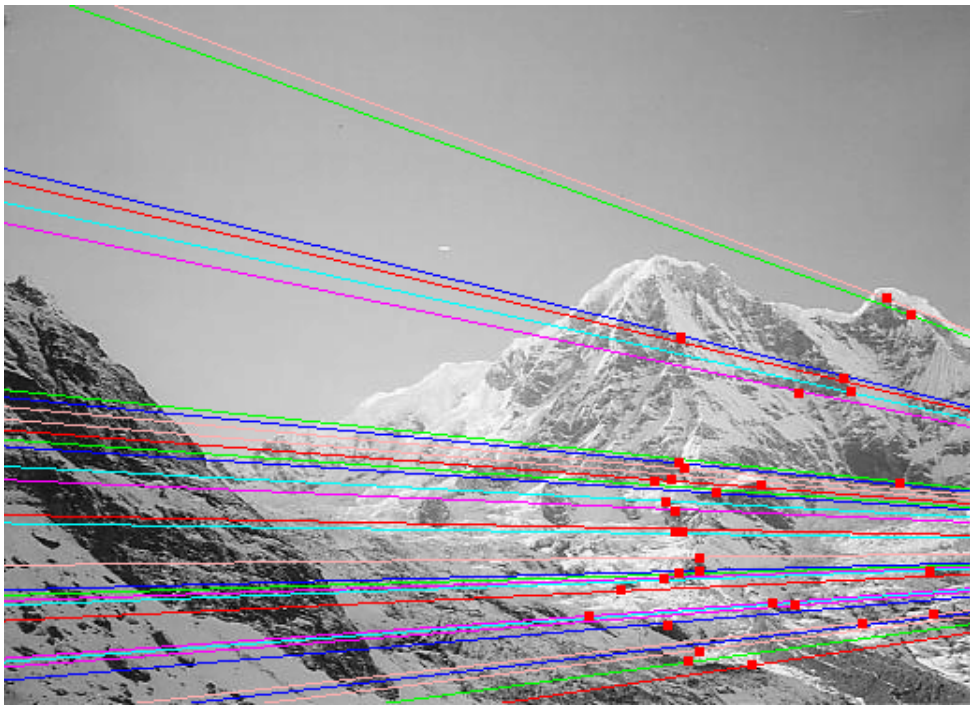


same camera objectives (nadirs), but
 different orientation for the 2 images
 (that is, rotated around the same nadir)

Point Correspondence

Need to create list of matched points between the images in order to solve for the epipolar projection

manually making a point list from the corners from the edge extractor used with “outdoor mode”:




stereo projection fit
to 32 points already
known to match
shows what the
epipolar projections
should be when the
corner find + corner
match + stereo
projection solve are
correctly
automated.


nMatched=32
avgDist=0.281
stDev=0.508

The Brown & Lowe 2003 images: point matching difficult because image intersection << difference

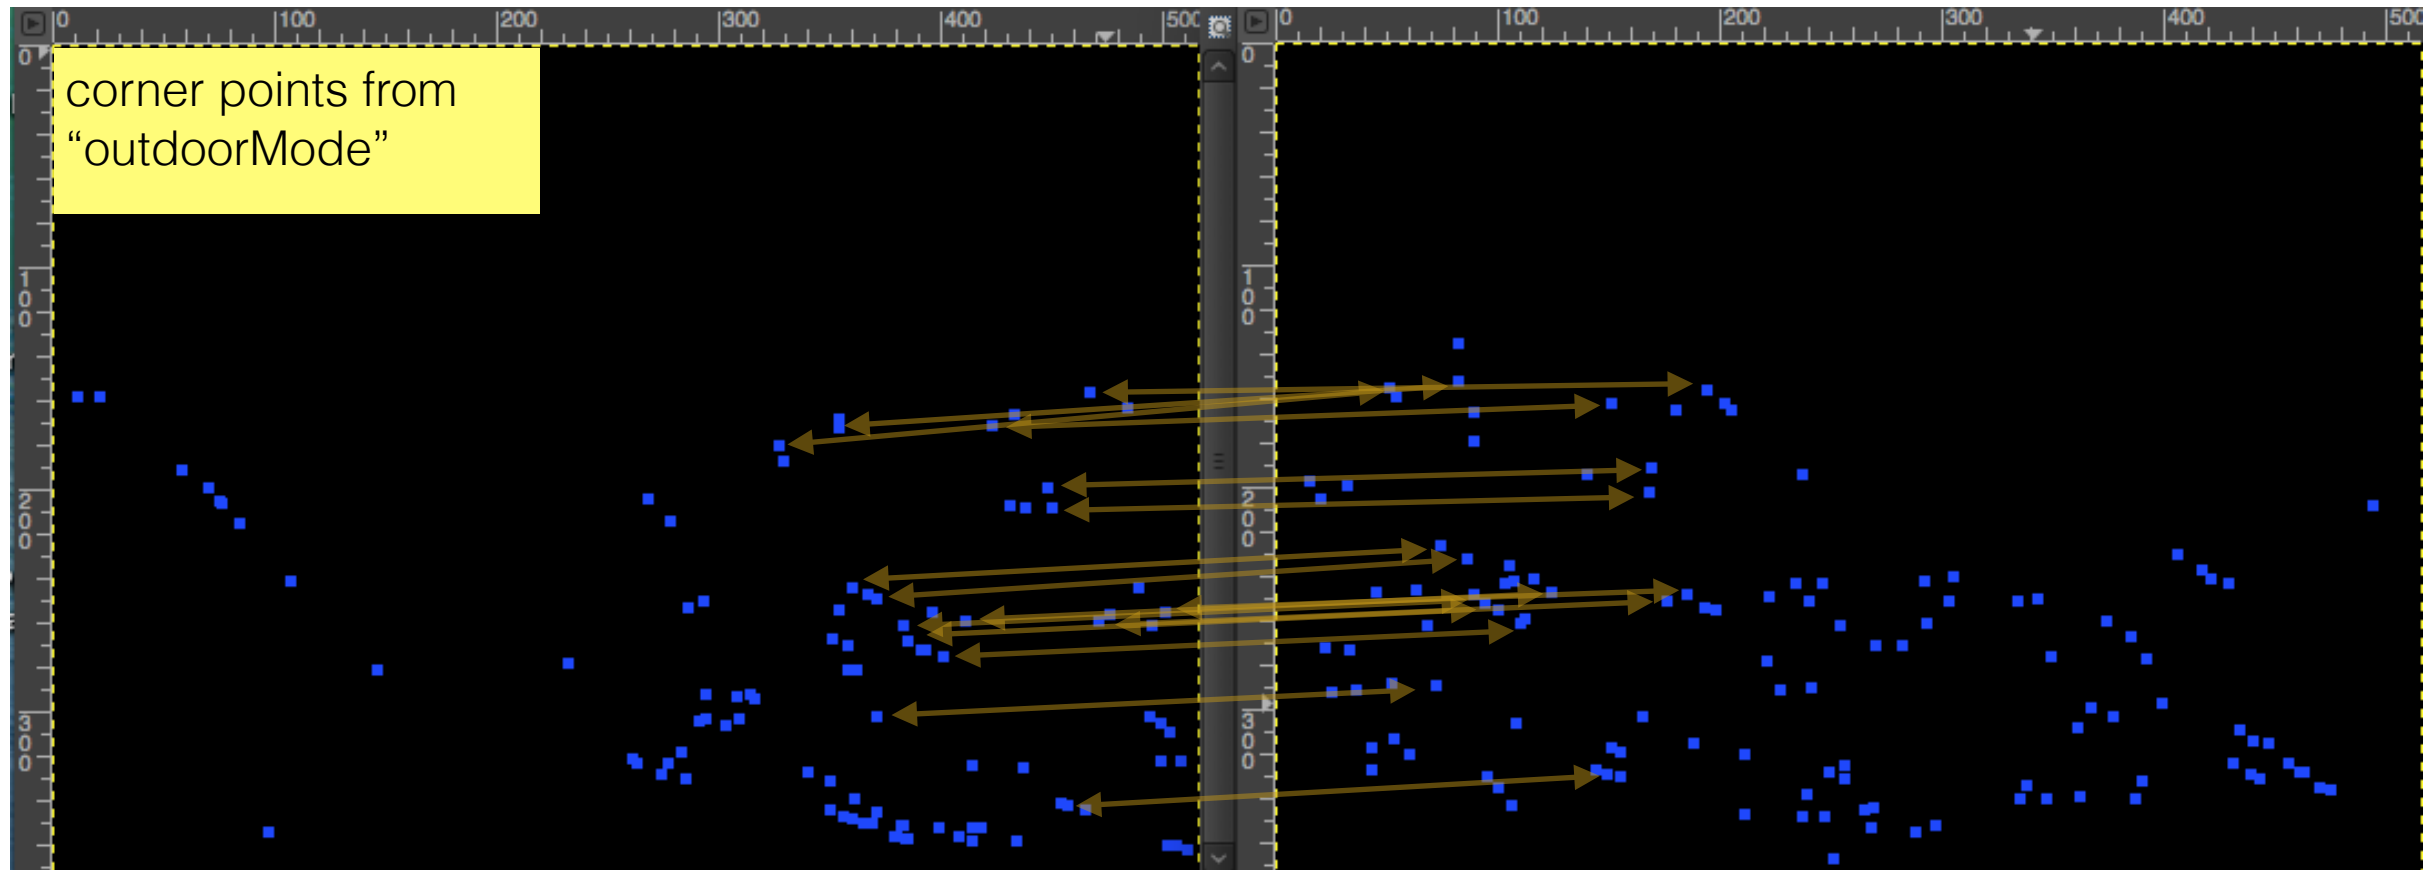
inflection points from
"outdoorMode"





could consider distinct curves and their immediate neighbors, but that would be many more points:



The Brown & Lowe 2003 images: point matching difficult because image intersection < difference

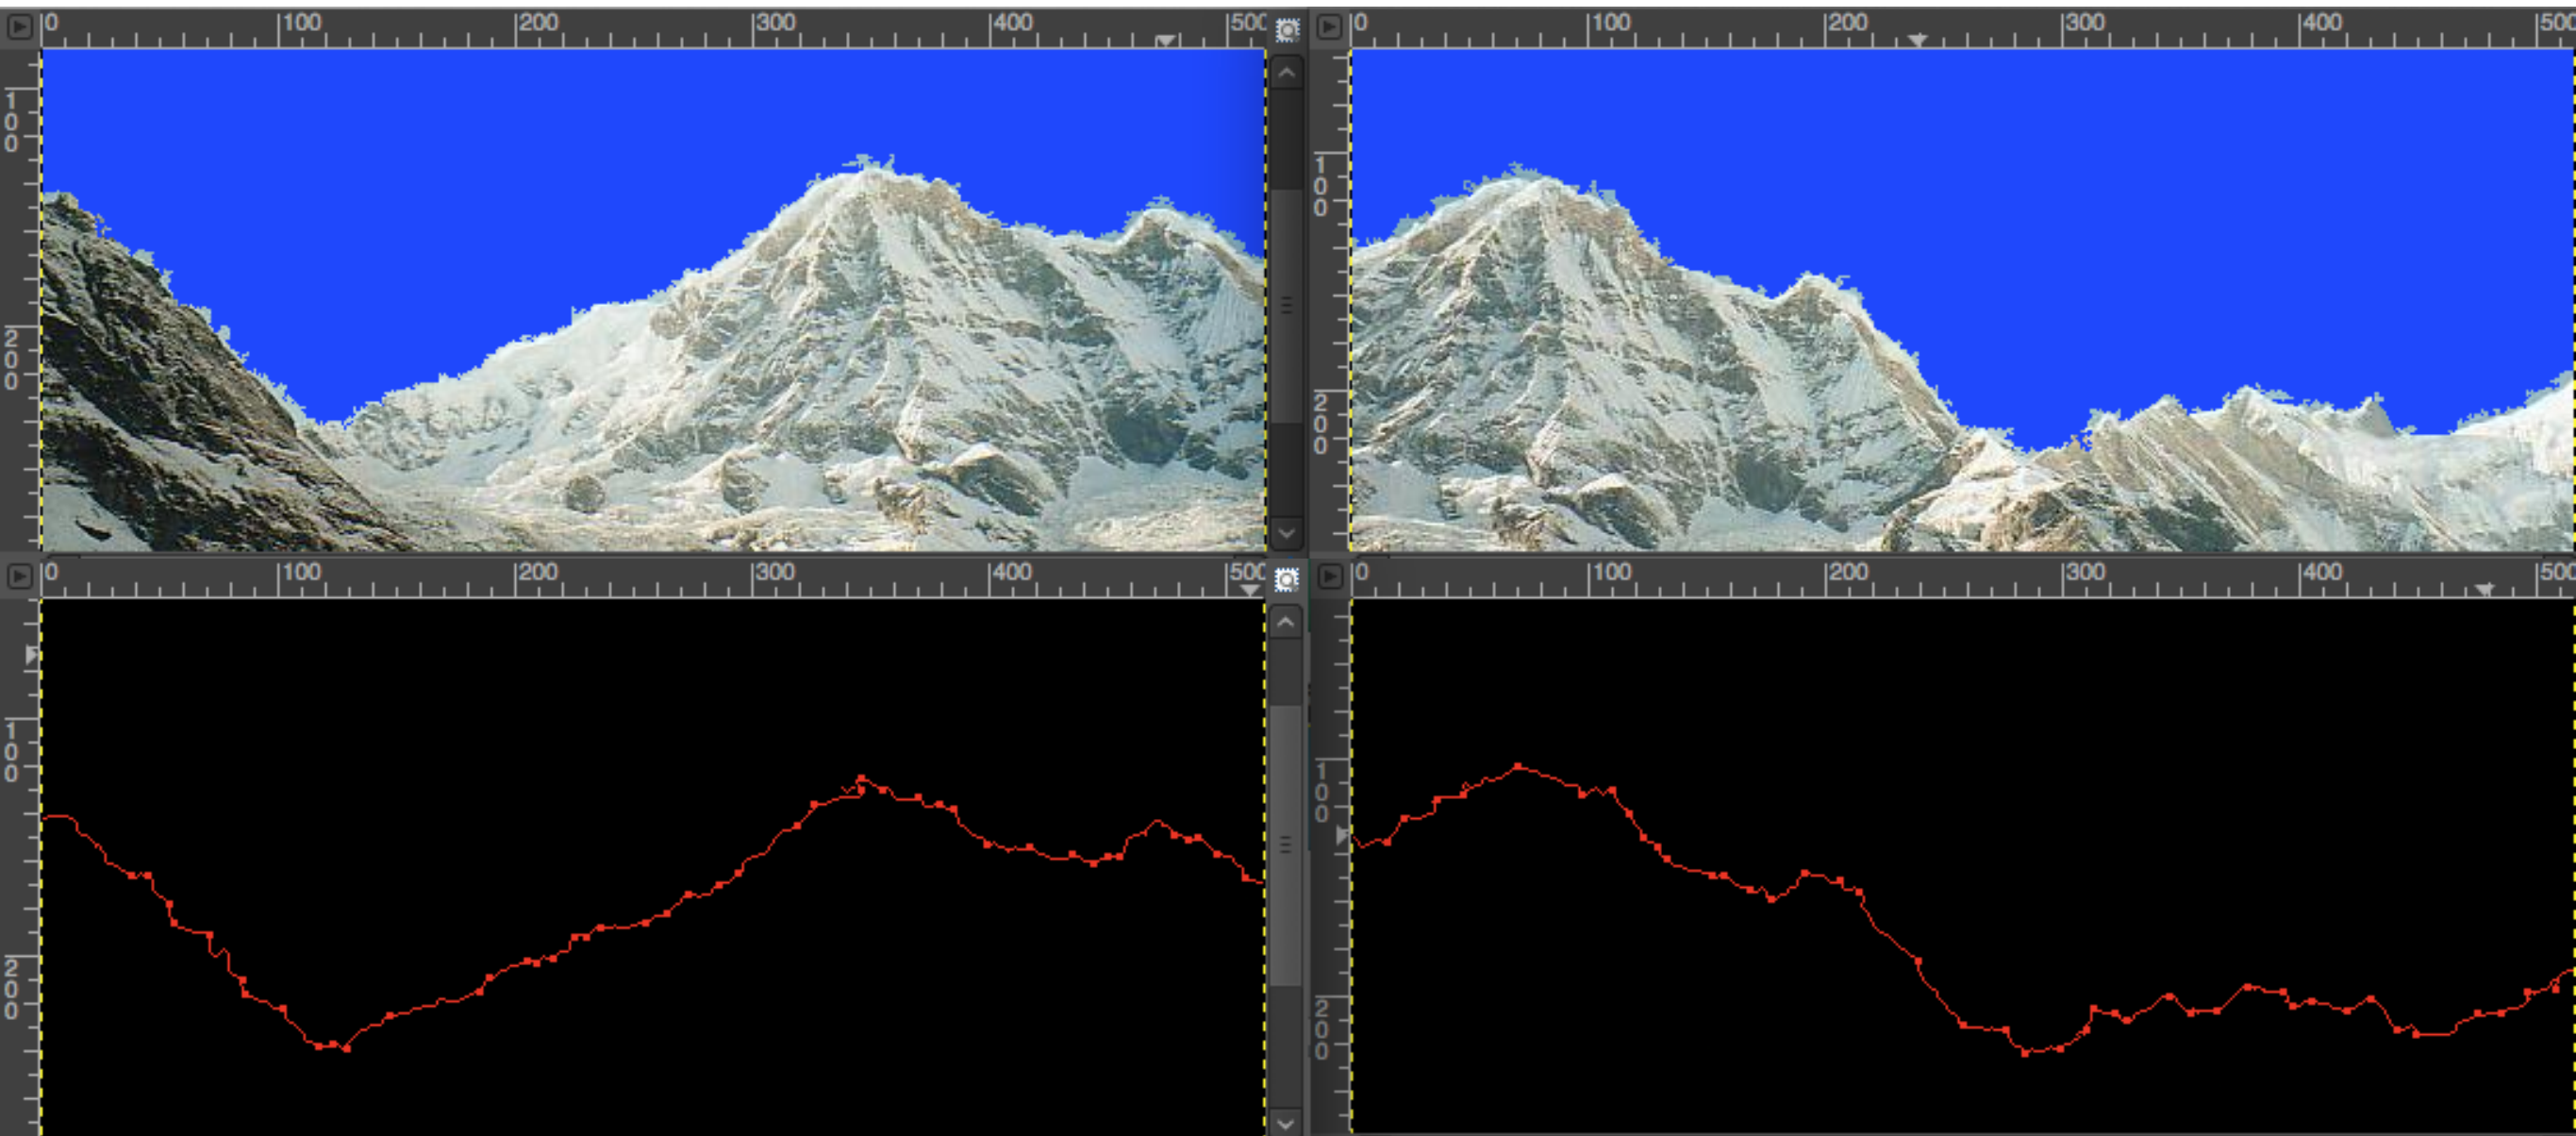


could consider distinct curves and their immediate neighbors, but that would be many more points:



For the outdoor images, can find the sky and create a sky mask and also create corners from just the skyline.
(see skyline_extraction.pdf)

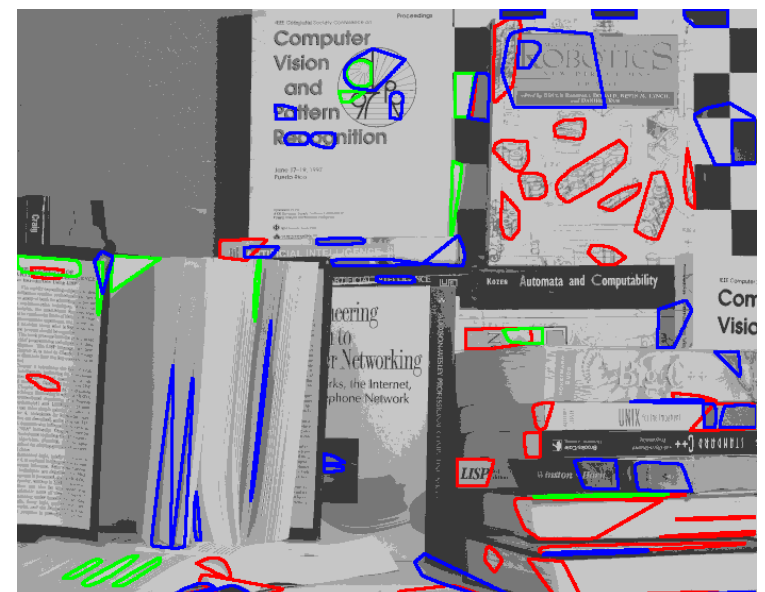


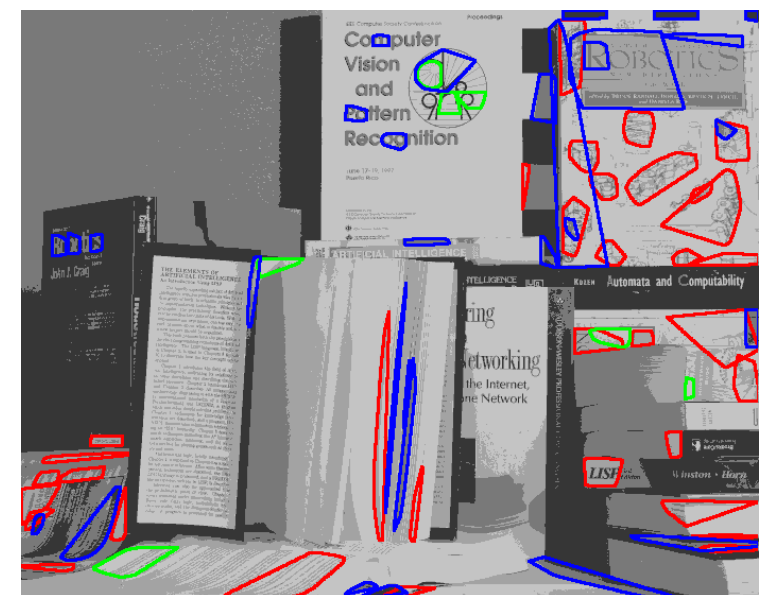


This sky mask helps to pre-process the image before feature finding and correspondence.



Trying blob feature creation and matching methods in next slides.

- perform histogram equalization if mean is too far from median or the two images are too different for those params.
- color segmentation of ***k=4*** to reduce image to 4 bands of intensities.
- contiguous pixel group finder for each of the 4 bands using point limits of: smallestGroupLimit = ***100***; largestGroupLimit = ***1000***;







Would prefer fewer features because the pairwise calculations of transformation are $(N1*(N1-1)/2) * (N2*(N2-1)/2)$ which is roughly N^4 . Note that trying combinations of rotation and translation (and scale) increase with the image sizes but are very large numbers too for reasonable accuracy.



Even better:

- perform histogram equalization if mean is too far from median or the two images are too different for those params.
- color segmentation of **$k=2$** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of: smallestGroupLimit = **100**; largestGroupLimit = **1000**;

This is a smaller number of features that almost has all of the major blobs one would want to use for matching.



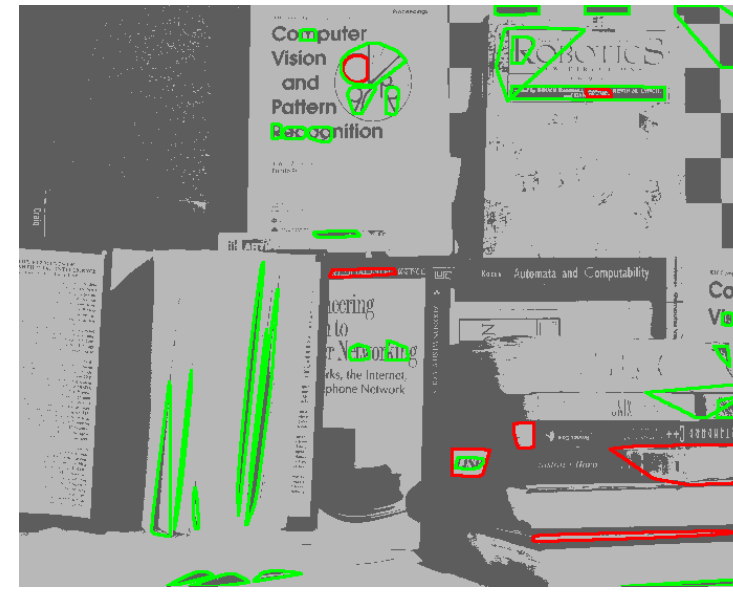







Even better:

- perform histogram equalization if mean is too far from median or the two images are too different for those stats.
- color segmentation of **$k=2$** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of: smallestGroupLimit = **100**; largestGroupLimit = **5000**;

correspondence looks like enough for an initial Euclidean solution

NOTE that blobs as first transformation solution works better for images like the middle where in contrast if done with corners, the ridge line in the differences would match more strongly than the true corner matches.











After Euclidean transformation from blobs, need to make lists of matching points for input to epipolar projection solver.

- gaussian blur with $\sigma=2$
- scale space curvature corners

While finding the matches, may need to increase the tolerance or change the solution slightly across the image due to projection.






With the transformation already roughly solved, making point lists even for a large number of corners is fast.



Alternatively for the blob (feature) finding, can reduce the image size:

- perform histogram equalization if mean is too far from median or the two images are too different for those params.
- ***bin the image to $< 200 \times 200$***
- color segmentation of ***$k=2$*** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of:
smallestGroupLimit = ***$100/(\text{binFactor}^2)$*** ; largestGroupLimit = ***$5000/(\text{binFactor}^2)$*** ;

This solution doesn't have a better matchable to non-matchable blob fraction than the full image, so will not use the binned processing.










After Euclidean transformation from blobs, need to make lists of matching points for input to epipolar projection solver.

- No gaussian blur
- scale space curvature corners



Note that this stage looks like it it best to perform using the full image, even though solving the Euclidean with the binned features is still a good first stage.

The binned blob features do not look better than the full image size blobs, but here are the corners from them just to have followed it through.







corners from images pre-processed by segmentation and binning (not using a Gaussian pyramid yet, but will later).





contiguous regions within size ranges were found then convex hulls constructed around them.




corners filtered to those hulls to reduce the corner list to regions of interest.

the number of possible true matches is less than half of the total number of corners, but feature matching with color image descriptors should help match them.





corners filtered




need to determine a gradient orientation which is consistent w.r.t. the corner.

Trying with the full images first to make sure the method and math are correct, then will see if can derive the same answer using the binned images.



gradient orientation. making corners with CSS means already have useful corner information and edge information to determine this more precisely, so am using those here.





A vector perpendicular to k_{\max} on this edge can be used to calculate the orientation of this region (where the region is defined by being a CSS corner).


deriving the perpendicular angle from the points directly to the left and right of the maximum of curvature for points which have curvature and weighting the derived perpendicular angle using weights that are the fraction of the sqsum of their distances from the maximum of curvature.

k	x	y	
0.24	267	84	dx,dy=(1,1), perp=45, dist=math.sqrt(2)
0.26	268	85	<-max k
0.21	268	86	dx,dy=(0,1), perp=-45 dist=math.sqrt(1)

dTot=1+math.sqrt(2)=2.414
orientation = (math.sqrt(2)/2.414)*45 + (1/2.414)*(-45) = 7.7



k is curvature





maximum in curvature

extract curvature from the highest resolution scale space curves:

idx=164 (268.0,85.0)


k	x	y
0.08	for (258, 79)	
0.04	for (259, 79)	
0.01	for (260, 80)	
0.00	for (261, 80)	
0.02	for (262, 81)	
0.03	for (263, 82)	
0.06	for (264, 83)	
0.10	for (265, 83)	
0.17	for (266, 83)	
0.24	for (267, 84)	
0.26	for (268, 85) <-k_max	
0.21	for (268, 86)	
0.13	for (267, 87)	
0.04	for (266, 88)	
0.04	for (266, 89)	
0.07	for (266, 90)	
0.06	for (265, 90)	
0.00	for (266, 91)	
0.07	for (267, 92)	
0.13	for (268, 93)	
0.15	for (268, 94)	

gradient orientation.






extract the final curvature from the highest resolution
scale space curves:
idx=112 (143.0,255.0)

k[110]=0.01	for (143, 253)	
k[111]=0.13	for (143, 254)	
k[112]=0.34	for (143, 255)	dx,dy=(-1,0)
k[113]=0.43	for (142, 255)	← k_max
k[114]=0.35	for (141, 255)	dx,dy=(-1,0)
k[115]=0.28	for (140, 255)	
k[116]=0.24	for (139, 255)	






$$\text{orientation} = (1/2)*90 + (1/2)*90 = 90$$



corners filtered (not finished here, but keeping snapshots)



A quick look at color segmentation with $k=3$ followed by binary segmentation, binned to size $< 200 \times 200$, followed by corners worked fine for 2 image sets, but not the third so the other methods in previous slides are preferred (corners filtered to blobs, matching, calculating transformation, then creating correspondence with all corners).



For the stereoscopic above, there's more than translation so this can only be the start of the solution

